

sentiment-analysis-distilbert

November 28, 2024

1 Sentiment Analysis Using DistilBERT

1.1 Project Details

- **Name:** Shrinidhi Krpete Thimmegowda
- **Email:** krpetesh@msu.edu
- **Dataset:**
 - **Source:** Amazon Product Reviews
 - **Size:** 1000 reviews
 - **Class Distribution:**
 - * 500 Positive (Label: 1)
 - * 500 Negative (Label: 0)

```
[1]: import numpy as np
import pandas as pd
from transformers import DistilBertTokenizer, DistilBertModel, AdamW, \
    get_scheduler
from torch.utils.data import DataLoader, TensorDataset, random_split
import torch
from tqdm import tqdm
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
```

c:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\tqdm\auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See

https://ipywidgets.readthedocs.io/en/stable/user_install.html

```
from .autonotebook import tqdm as notebook_tqdm
```

Step 2: Load data Analyze

```
[ ]: with open('C:/Users/Admin/Downloads/archive (9)/sentiment labelled sentences/
    amazon_cells_labelled.txt', 'r') as file:
    # Read all the content of the file
    data = file.read()
```

```
[3]: # Split the data into lines
lines = data.strip().split("\n")
```

```

# Split each line into two parts: review and sentiment
rows = [line.rsplit("\t", 1) for line in lines]

# Create a DataFrame
df = pd.DataFrame(rows, columns=["Review", "Sentiment"])

# Convert the 'Sentiment' column to integers
df["Sentiment"] = df["Sentiment"].astype(int)

# Display the DataFrame
print(df)

```

	Review	Sentiment
0	So there is no way for me to plug it in here i...	0
1	Good case, Excellent value.	1
2	Great for the jawbone.	1
3	Tied to charger for conversations lasting more...	0
4	The mic is great.	1
..
995	The screen does get smudged easily because it ...	0
996	What a piece of junk.. I lose more calls on th...	0
997	Item Does Not Match Picture.	0
998	The only thing that disappoint me is the infra...	0
999	You can not answer calls with the unit, never ...	0

[1000 rows x 2 columns]

```
[4]: df.dtypes
```

```

[4]: Review      object
     Sentiment    int32
     dtype: object

```

```
[5]: df['Sentiment'].unique()
```

```
[5]: array([0, 1])
```

```

[6]: texts = list(df["Review"]) # Reviews
     labels = list(df["Sentiment"]) # Sentiments (0 or 1)

```

Step 2: Tokenize Data

```

[7]: tokenizer = DistilBertTokenizer.from_pretrained("distilbert-base-uncased")
     inputs = tokenizer(
         texts,
         padding=True,

```

```

        truncation=True,
        max_length=128,
        return_tensors="pt"
    )
    labels = torch.tensor(labels)

```

Step 3: Create Dataset and DataLoader

```

[8]: dataset = TensorDataset(inputs["input_ids"], inputs["attention_mask"], labels)
    train_size = int(0.8 * len(dataset))
    val_size = len(dataset) - train_size
    train_dataset, val_dataset = random_split(dataset, [train_size, val_size])

    train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
    val_loader = DataLoader(val_dataset, batch_size=16)

```

Step 4: Initialize DistilBERT

```

[9]: distilbert_model = DistilBertModel.from_pretrained("distilbert-base-uncased")
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    distilbert_model.to(device)

```

```

[9]: DistilBertModel(
  (embeddings): Embeddings(
    (word_embeddings): Embedding(30522, 768, padding_idx=0)
    (position_embeddings): Embedding(512, 768)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
  (transformer): Transformer(
    (layer): ModuleList(
      (0-5): 6 x TransformerBlock(
        (attention): DistilBertSdpaAttention(
          (dropout): Dropout(p=0.1, inplace=False)
          (q_lin): Linear(in_features=768, out_features=768, bias=True)
          (k_lin): Linear(in_features=768, out_features=768, bias=True)
          (v_lin): Linear(in_features=768, out_features=768, bias=True)
          (out_lin): Linear(in_features=768, out_features=768, bias=True)
        )
        (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (ffn): FFN(
          (dropout): Dropout(p=0.1, inplace=False)
          (lin1): Linear(in_features=768, out_features=3072, bias=True)
          (lin2): Linear(in_features=3072, out_features=768, bias=True)
          (activation): GELUActivation()
        )
        (output_layer_norm): LayerNorm((768,), eps=1e-12,
        elementwise_affine=True)
      )
    )
  )
)

```

```
)
)
)
)
```

Step 5: Fine-Tune DistilBERT

```
[10]: optimizer = AdamW(distilbert_model.parameters(), lr=2e-5)
num_training_steps = len(train_loader) * 3 # 3 epochs
scheduler = get_scheduler("linear", optimizer=optimizer, num_warmup_steps=0,
    ↪ num_training_steps=num_training_steps)

distilbert_model.train()
for epoch in range(3): # Adjust number of epochs as needed
    progress_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}")
    for batch in progress_bar:
        input_ids, attention_mask, _ = [b.to(device) for b in batch]

        # Forward pass
        outputs = distilbert_model(input_ids, attention_mask=attention_mask)
        embeddings = outputs.last_hidden_state[:, 0, :] # First token
        ↪ ([CLS]-like embedding)

        # Compute dummy loss (optional for fine-tuning purposes)
        loss = embeddings.norm(2) # Example dummy loss
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        scheduler.step()

distilbert_model.eval()
```

c:\Users\Admin\AppData\Local\Programs\Python\Python312\Lib\site-packages\transformers\optimization.py:591: FutureWarning: This implementation of AdamW is deprecated and will be removed in a future version. Use the PyTorch implementation torch.optim.AdamW instead, or set `no_deprecation_warning=True` to disable this warning

```
warnings.warn(
Epoch 1: 100%|      | 50/50 [02:26<00:00,  2.93s/it]
Epoch 2: 100%|      | 50/50 [02:09<00:00,  2.60s/it]
Epoch 3: 100%|      | 50/50 [02:04<00:00,  2.49s/it]
```

```
[10]: DistilBertModel(
  (embeddings): Embeddings(
    (word_embeddings): Embedding(30522, 768, padding_idx=0)
    (position_embeddings): Embedding(512, 768)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
```

```

        (dropout): Dropout(p=0.1, inplace=False)
    )
    (transformer): Transformer(
      (layer): ModuleList(
        (0-5): 6 x TransformerBlock(
          (attention): DistilBertSdpaAttention(
            (dropout): Dropout(p=0.1, inplace=False)
            (q_lin): Linear(in_features=768, out_features=768, bias=True)
            (k_lin): Linear(in_features=768, out_features=768, bias=True)
            (v_lin): Linear(in_features=768, out_features=768, bias=True)
            (out_lin): Linear(in_features=768, out_features=768, bias=True)
          )
          (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (ffn): FFN(
            (dropout): Dropout(p=0.1, inplace=False)
            (lin1): Linear(in_features=768, out_features=3072, bias=True)
            (lin2): Linear(in_features=3072, out_features=768, bias=True)
            (activation): GELUActivation()
          )
          (output_layer_norm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
        )
      )
    )
  )
)

```

Step 6: Extract Fine-Tuned Embeddings

```

[11]: with torch.no_grad():
        outputs = distilbert_model(**inputs)
        embeddings = outputs.last_hidden_state[:, 0, :].cpu().numpy() # Extract
        ↪ [CLS]-like embeddings

```

Step 7: Train Logistic Regression

```

[12]: # Split data into train and validation sets
X_train, X_val, y_train, y_val = train_test_split(embeddings, labels.numpy(),
        ↪ test_size=0.2, random_state=42)

# Train logistic regression
logistic_model = LogisticRegression(max_iter=1000)
logistic_model.fit(X_train, y_train)

```

```

[12]: LogisticRegression(max_iter=1000)

```

Step 8: Evaluate Logistic Regression

```
[13]: y_pred = logistic_model.predict(X_val)
accuracy = accuracy_score(y_val, y_pred)
print(f"Validation Accuracy: {accuracy:.4f}")
print("Classification Report:")
print(classification_report(y_val, y_pred, target_names=["Negative", "Positive"]))
```

Validation Accuracy: 0.8900

Classification Report:

	precision	recall	f1-score	support
Negative	0.87	0.90	0.88	93
Positive	0.91	0.88	0.90	107
accuracy			0.89	200
macro avg	0.89	0.89	0.89	200
weighted avg	0.89	0.89	0.89	200

1.2 Results Summary

1.2.1 Validation Metrics

- **Validation Accuracy: 89.0%**
 - The model correctly classified 89% of the validation dataset.

1.2.2 Class-Specific Metrics

- **Negative Class (Label: 0):**
 - **Precision:** 87% – Out of all reviews predicted as **Negative**, 87% were actually **Negative**.
 - **Recall:** 90% – Out of all actual **Negative** reviews, 90% were correctly identified.
 - **F1-Score:** 88% – A balanced performance metric combining precision and recall.
- **Positive Class (Label: 1):**
 - **Precision:** 91% – Out of all reviews predicted as **Positive**, 91% were actually **Positive**.
 - **Recall:** 88% – Out of all actual **Positive** reviews, 88% were correctly identified.
 - **F1-Score:** 90% – Strong overall performance in identifying **Positive** reviews.

1.2.3 Overall Metrics

- **Macro Average:**
 - Precision: 89%
 - Recall: 89%
 - F1-Score: 89%
- **Weighted Average:**
 - Precision: 89%
 - Recall: 89%
 - F1-Score: 89%

1.3 Key Takeaways

- The model achieves a strong balance between precision, recall, and F1-score across both classes.
- **Positive reviews** are identified with slightly higher precision (91%), reducing false positives.
- **Negative reviews** achieve slightly higher recall (90%), reducing false negatives.
- The balanced metrics indicate the model effectively handles both classes with consistent performance.